

AUS_Labour_Market_Tech_Sector_Analysis_by_Dohyeon_Lee.R

dohye

2022-01-17

```
# Description: Analysis of the current state of the Australian tech labour market
#               and predicting the future trend.
# Part A: Import Data
# Part B: Data Preparation and Wrangling
# Part C: Initial Analysis
# Part D: Data Visualization
# Part E: Exploratory Data Analysis, Data-Driven Modeling
# Part F: Evaluate Results
# Part G: Forecast trend

##### Part A #####
# install packages if it is not installed
list.of.packages <- c("tidyverse", "stringr", "reshape2", "Metrics", "scales", "lubridate",
                      "RSocrata", "zoo", "xts", "httr", "xlsx", "readxl", "sjmisc", "dplyr",
                      "rpart", "rattle", "randomForest", "gsubfn", "e1071")
new.packages <- list.of.packages[!(list.of.packages %in% installed.packages()[, "Package"])]
if(length(new.packages)) install.packages(new.packages)

# Load packages installed
lapply(list.of.packages, library, character.only = TRUE)

## -- Attaching packages ----- tidyverse 1.3.1 --
## v ggplot2 3.3.5     v purrr    0.3.4
## v tibble   3.1.3     v dplyr    1.0.7
## v tidyr    1.1.3     v stringr  1.4.0
## v readr    2.0.1     vforcats  0.5.1

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()   masks stats::lag()

##
## Attaching package: 'reshape2'

## The following object is masked from 'package:tidyverse':
##   smiths
```

```

## 
## Attaching package: 'scales'

## The following object is masked from 'package:purrr':
## 
##     discard

## The following object is masked from 'package:readr':
## 
##     col_factor

## 
## Attaching package: 'lubridate'

## The following objects are masked from 'package:base':
## 
##     date, intersect, setdiff, union

## Warning: package 'RSocrata' was built under R version 4.1.2

## 
## Attaching package: 'zoo'

## The following objects are masked from 'package:base':
## 
##     as.Date, as.Date.numeric

## 
## Attaching package: 'xts'

## The following objects are masked from 'package:dplyr':
## 
##     first, last

## Warning: package 'xlsx' was built under R version 4.1.2

## java.home option:

## JAVA_HOME environment variable: C:\Program Files (x86)\Java\jdk1.6.0_23

## Warning in fun(libname, pkgname): Java home setting is INVALID, it will be ignored.
## Please do NOT set it unless you want to override system settings.

## Warning: package 'sjmisc' was built under R version 4.1.2

## 
## Attaching package: 'sjmisc'

## The following object is masked from 'package:purrr':
## 
##     is_empty

```

```

## The following object is masked from 'package:tidyr':
##
##     replace_na

## The following object is masked from 'package:tibble':
##
##     add_case

## Loading required package: bitops

## Rattle: A free graphical interface for data science with R.
## Version 5.4.0 Copyright (c) 2006-2020 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.

## randomForest 4.6-14

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'

## The following object is masked from 'package:rattle':
##
##     importance

## The following object is masked from 'package:dplyr':
##
##     combine

## The following object is masked from 'package:ggplot2':
##
##     margin

## Warning: package 'gsubfn' was built under R version 4.1.2

## Loading required package: proto

## Warning: package 'proto' was built under R version 4.1.2

## [[1]]
## [1] "forcats"    "stringr"     "dplyr"       "purrr"       "readr"       "tidyr"
## [7] "tibble"      "ggplot2"     "tidyverse"   "stats"       "graphics"   "grDevices"
## [13] "utils"       "datasets"    "methods"     "base"
##
## [[2]]
## [1] "forcats"    "stringr"     "dplyr"       "purrr"       "readr"       "tidyr"
## [7] "tibble"      "ggplot2"     "tidyverse"   "stats"       "graphics"   "grDevices"
## [13] "utils"       "datasets"    "methods"     "base"
##
## [[3]]
## [1] "reshape2"    "forcats"     "stringr"     "dplyr"       "purrr"       "readr"

```

```

## [7] "tidyverse" "stats"       "graphics"
## [13] "grDevices" "utils"       "datasets"    "methods"     "base"
##
## [[4]]
## [1] "Metrics"   "reshape2"   "forcats"    "stringr"    "dplyr"      "purrr"
## [7] "readr"      "tidyverse"  "tibble"     "ggplot2"    "tidyverse"  "stats"
## [13] "graphics"   "grDevices"   "utils"      "datasets"   "methods"    "base"
##
## [[5]]
## [1] "scales"    "Metrics"    "reshape2"   "forcats"   "stringr"    "dplyr"
## [7] "purrr"     "readr"      "tidyverse"  "tibble"    "ggplot2"    "tidyverse"
## [13] "stats"     "graphics"   "grDevices"  "utils"     "datasets"   "methods"
## [19] "base"
##
## [[6]]
## [1] "lubridate" "scales"    "Metrics"    "reshape2"   "forcats"   "stringr"
## [7] "dplyr"      "purrr"     "readr"      "tidyverse"  "tibble"    "ggplot2"
## [13] "tidyverse"  "stats"     "graphics"   "grDevices"  "utils"     "datasets"
## [19] "methods"   "base"
##
## [[7]]
## [1] "RSocrata"  "lubridate"  "scales"    "Metrics"    "reshape2"   "forcats"
## [7] "stringr"   "dplyr"     "purrr"     "readr"     "tidyverse"  "tibble"
## [13] "ggplot2"   "tidyverse"  "stats"     "graphics"   "grDevices"  "utils"
## [19] "datasets"  "methods"   "base"
##
## [[8]]
## [1] "zoo"        "RSocrata"  "lubridate"  "scales"    "Metrics"    "reshape2"
## [7] "forcats"   "stringr"   "dplyr"     "purrr"     "readr"     "tidyverse"
## [13] "tibble"    "ggplot2"   "tidyverse"  "stats"     "graphics"   "grDevices"
## [19] "utils"     "datasets"  "methods"   "base"
##
## [[9]]
## [1] "xts"        "zoo"       "RSocrata"  "lubridate"  "scales"    "Metrics"
## [7] "reshape2"   "forcats"   "stringr"   "dplyr"     "purrr"     "readr"
## [13] "tidyverse"  "tibble"    "ggplot2"   "tidyverse"  "stats"     "graphics"
## [19] "grDevices"  "utils"     "datasets"  "methods"   "base"
##
## [[10]]
## [1] "httr"       "xts"       "zoo"       "RSocrata"  "lubridate"  "scales"
## [7] "Metrics"    "reshape2"   "forcats"   "stringr"   "dplyr"     "purrr"
## [13] "readr"      "tidyverse"  "tibble"    "ggplot2"   "tidyverse"  "stats"
## [19] "graphics"   "grDevices"  "utils"     "datasets"  "methods"   "base"
##
## [[11]]
## [1] "xlsx"       "httr"      "xts"       "zoo"       "RSocrata"  "lubridate"
## [7] "scales"     "Metrics"   "reshape2"   "forcats"   "stringr"   "dplyr"
## [13] "purrr"     "readr"     "tidyverse"  "tibble"    "ggplot2"   "tidyverse"
## [19] "stats"     "graphics"   "grDevices"  "utils"     "datasets"  "methods"
## [25] "base"
##
## [[12]]
## [1] "readxl"     "xlsx"      "httr"      "xts"       "zoo"       "RSocrata"
## [7] "lubridate"  "scales"   "Metrics"   "reshape2"   "forcats"   "stringr"

```

```

## [13] "dplyr"      "purrr"       "readr"       "tidyverse"   "tibble"      "ggplot2"
## [19] "tidyverse"   "stats"        "graphics"    "grDevices"   "utils"       "datasets"
## [25] "methods"     "base"         " "
## 
## [[13]]
## [1] "sjmisc"     "readxl"      "xlsx"       "httr"        "xts"        "zoo"
## [7] "RSocrata"   "lubridate"   "scales"     "Metrics"     "reshape2"   "forcats"
## [13] "stringr"    "dplyr"       "purrr"     "readr"       "tidyverse"  "tibble"
## [19] "ggplot2"    "tidyverse"   "stats"      "graphics"   "grDevices"  "utils"
## [25] "datasets"   "methods"     "base"       " "
## 
## [[14]]
## [1] "sjmisc"     "readxl"      "xlsx"       "httr"        "xts"        "zoo"
## [7] "RSocrata"   "lubridate"   "scales"     "Metrics"     "reshape2"   "forcats"
## [13] "stringr"    "dplyr"       "purrr"     "readr"       "tidyverse"  "tibble"
## [19] "ggplot2"    "tidyverse"   "stats"      "graphics"   "grDevices"  "utils"
## [25] "datasets"   "methods"     "base"       " "
## 
## [[15]]
## [1] "rpart"       "sjmisc"      "readxl"      "xlsx"       "httr"        "xts"
## [7] "zoo"          "RSocrata"   "lubridate"   "scales"     "Metrics"     "reshape2"
## [13] "forcats"     "stringr"     "dplyr"       "purrr"     "readr"       "tidyverse"
## [19] "tibble"       "ggplot2"     "tidyverse"   "stats"      "graphics"   "grDevices"
## [25] "utils"       "datasets"    "methods"     "base"       " "
## 
## [[16]]
## [1] "rattle"      "bitops"      "rpart"      "sjmisc"     "readxl"      "xlsx"
## [7] "httr"        "xts"        "zoo"        "RSocrata"   "lubridate"   "scales"
## [13] "Metrics"     "reshape2"    "forcats"    "stringr"    "dplyr"       "purrr"
## [19] "readr"        "tidyverse"   "tibble"     "ggplot2"    "tidyverse"   "stats"
## [25] "graphics"    "grDevices"   "utils"      "datasets"   "methods"     "base"
## 
## [[17]]
## [1] "randomForest" "rattle"      "bitops"      "rpart"      "sjmisc"
## [6] "readxl"       "xlsx"       "httr"        "xts"        "zoo"
## [11] "RSocrata"    "lubridate"   "scales"     "Metrics"     "reshape2"
## [16] "forcats"     "stringr"     "dplyr"       "purrr"     "readr"
## [21] "tidyverse"   "tibble"     "ggplot2"    "tidyverse"   "stats"
## [26] "graphics"    "grDevices"   "utils"      "datasets"   "methods"
## [31] "base"         " "
## 
## [[18]]
## [1] "gsubfn"      "proto"       "randomForest" "rattle"     "bitops"
## [6] "rpart"        "sjmisc"     "readxl"      "xlsx"       "httr"
## [11] "xts"          "zoo"        "RSocrata"   "lubridate"   "scales"
## [16] "Metrics"      "reshape2"    "forcats"    "stringr"    "dplyr"
## [21] "purrr"        "readr"       "tidyverse"   "tibble"     "ggplot2"
## [26] "tidyverse"   "stats"       "graphics"   "grDevices"   "utils"
## [31] "datasets"    "methods"     "base"       " "
## 
## [[19]]
## [1] "e1071"        "gsubfn"     "proto"       "randomForest" "rattle"
## [6] "bitops"        "rpart"      "sjmisc"     "readxl"      "xlsx"
## [11] "httr"          "xts"        "zoo"        "RSocrata"   "lubridate"

```

```

## [16] "scales"      "Metrics"       "reshape2"      "forcats"       "stringr"
## [21] "dplyr"        "purrr"         "readr"         "tidyR"         "tibble"
## [26] "ggplot2"      "tidyverse"     "stats"         "graphics"     "grDevices"
## [31] "utils"         "datasets"      "methods"       "base"

# Import list of time series file names
data_ts_list = list.files(path='data/', pattern="EQ", full.names=TRUE)

# Load and trim names of data frames
data_name_list <- NULL #initialise list name for list of data frames
data_name_list_trimed <- NULL
for (i in 1:length(data_ts_list)){ # extract titles of each data frames and trim the strings into more
  data_name_list[i] <- gsub( ".*Employed (.+) job .*", "\\\1",
                            as.character(colnames(
                              read_excel(
                                data_ts_list[i], sheet = 3, range = cell_rows(2:2)))
                            )
                          )
  data_name_list_trimed[i] <- gsub(" ", "_", paste0("Employed ", data_name_list[i], " job"))
}

# Import time series files and store as tibble
for (i in 1:length(data_ts_list)){
  data_imported <- as_tibble(read_excel(data_ts_list[i],
                                         skip = 3,
                                         sheet = 3
                                         ), trimws("both"))
  names(data_imported)[1] <- "Date"
  names(data_imported)[grep('v', names(data_imported))] <- "Occupation"
  names(data_imported) <- gsub(" ", "_", names(data_imported))
  data_imported$Occupation <- gsub("[[:digit:]]", "", data_imported$Occupation)
  assign(paste(data_name_list_trimed[i]), data_imported) # rename each data frame imported to start with
}

#####
# Merge data frames into a master data frame
master_df <- get(data_name_list_trimed[1]) #initialise empty tibble data
for (i in 2:length(data_name_list_trimed)){ #merge every tibble data imported; data
  master_df <- merge(master_df, get(data_name_list_trimed[2]),
                      by = c("Date", "Occupation"), all.x = TRUE)
}

## Warning in merge.data.frame(master_df, get(data_name_list_trimed[2]), by =
## c("Date", : column names 'Sex.x', 'Employed_full-time_(000).x', 'Employed_part-
## time_(000).x', 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time_)_(000_Hours).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time_)_(000_Hours).x', 'Sex.y', 'Employed_full-time_(000).y', 'Employed_part-
## time_(000).y', 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time_)_(000_Hours).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time_)_(000_Hours).y' are duplicated in the result

## Warning in merge.data.frame(master_df, get(data_name_list_trimed[2]), by =

```

```

## c("Date", : column names 'Sex.x', 'Employed_full-time_(000).x', 'Employed_part-
## time_(000).x', 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)_(000_Hours).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)_(000_Hours).x', 'Sex.y', 'State_and_territory_(STT):_ASGS_(2011).x',
## 'Employed_full-time_(000).y', 'Employed_part-time_(000).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)_(000_Hours).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)_(000_Hours).y', 'State_and_territory_(STT):_ASGS_(2011).y' are duplicated
## in the result

## Warning in merge.data.frame(master_df, get(data_name_list_trimed[2]), by =
## c("Date", : column names 'Sex.x', 'Employed_full-time_(000).x', 'Employed_part-
## time_(000).x', 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)_(000_Hours).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)_(000_Hours).x', 'Sex.y', 'State_and_territory_(STT):_ASGS_(2011).x',
## 'Employed_full-time_(000).y', 'Employed_part-time_(000).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)_(000_Hours).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)_(000_Hours).y', 'Sex.x', 'State_and_territory_(STT):_ASGS_(2011).y',
## 'Employed_full-time_(000).x', 'Employed_part-time_(000).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)_(000_Hours).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)_(000_Hours).x', 'Sex.y', 'Employed_full-time_(000).y', 'Employed_part-
## time_(000).y', 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)_(000_Hours).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)_(000_Hours).y', 'State_and_territory_(STT):_ASGS_(2011).y' are duplicated in the result

## Warning in merge.data.frame(master_df, get(data_name_list_trimed[2]), by =
## c("Date", : column names 'Sex.x', 'Employed_full-time_(000).x', 'Employed_part-
## time_(000).x', 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)_(000_Hours).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)_(000_Hours).x', 'Sex.y', 'State_and_territory_(STT):_ASGS_(2011).x',
## 'Employed_full-time_(000).y', 'Employed_part-time_(000).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)_(000_Hours).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)_(000_Hours).y', 'Sex.x', 'State_and_territory_(STT):_ASGS_(2011).y',
## 'Employed_full-time_(000).x', 'Employed_part-time_(000).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)_(000_Hours).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)_(000_Hours).x', 'Sex.y', 'State_and_territory_(STT):_ASGS_(2011).x',
## 'Employed_full-time_(000).y', 'Employed_part-time_(000).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)_(000_Hours).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)_(000_Hours).y', 'State_and_territory_(STT):_ASGS_(2011).y' are duplicated

```



```

## 'Employed_full-time_(000).x', 'Employed_part-time_(000).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)__(000_Hours).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)__(000_Hours).x', 'Sex.y', 'State_and_territory_(STT):_ASGS_(2011).x',
## 'Employed_full-time_(000).y', 'Employed_part-time_(000).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)__(000_Hours).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)__(000_Hours).y', 'State_and_territory_(STT):_ASGS_(2011).y' are duplicated
## in the result

## Warning in merge.data.frame(master_df, get(data_name_list_trimed[2]), by =
## c("Date", : column names 'Sex.x', 'Employed_full-time_(000).x', 'Employed_part-
## time_(000).x', 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)__(000_Hours).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)__(000_Hours).x', 'Sex.y', 'State_and_territory_(STT):_ASGS_(2011).x',
## 'Employed_full-time_(000).y', 'Employed_part-time_(000).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)__(000_Hours).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)__(000_Hours).y', 'Sex.x', 'State_and_territory_(STT):_ASGS_(2011).y',
## 'Employed_full-time_(000).x', 'Employed_part-time_(000).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)__(000_Hours).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)__(000_Hours).x', 'Sex.y', 'State_and_territory_(STT):_ASGS_(2011).x',
## 'Employed_full-time_(000).y', 'Employed_part-time_(000).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)__(000_Hours).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)__(000_Hours).y', 'Sex.x', 'State_and_territory_(STT):_ASGS_(2011).y',
## 'Employed_full-time_(000).x', 'Employed_part-time_(000).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)__(000_Hours).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)__(000_Hours).x', 'Sex.y', 'State_and_territory_(STT):_ASGS_(2011).x',
## 'Employed_full-time_(000).x', 'Employed_part-time_(000).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)__(000_Hours).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)__(000_Hours).x', 'Sex.y', 'Employed_full-time_(000).y', 'Employed_part-
## time_(000).y', 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)__(000_Hours).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)__(000_Hours).y', 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)__(000_Hours).y' are duplicated in the result

## Warning in merge.data.frame(master_df, get(data_name_list_trimed[2]), by =

```

```

## c("Date", : column names 'Sex.x', 'Employed_full-time_(000).x', 'Employed_part-
## time_(000).x', 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)_('000_Hours).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)_('000_Hours).x', 'Sex.y', 'State_and_territory_(STT):_ASGS_(2011).x',
## 'Employed_full-time_(000).y', 'Employed_part-time_(000).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)_('000_Hours).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)_('000_Hours).y', 'Sex.x', 'State_and_territory_(STT):_ASGS_(2011).y',
## 'Employed_full-time_(000).x', 'Employed_part-time_(000).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)_('000_Hours).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)_('000_Hours).x', 'Sex.y', 'State_and_territory_(STT):_ASGS_(2011).x',
## 'Employed_full-time_(000).y', 'Employed_part-time_(000).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)_('000_Hours).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)_('000_Hours).y', 'Sex.x', 'State_and_territory_(STT):_ASGS_(2011).y',
## 'Employed_full-time_(000).x', 'Employed_part-time_(000).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)_('000_Hours).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)_('000_Hours).x', 'Sex.y', 'State_and_territory_(STT):_ASGS_(2011).x',
## 'Employed_full-time_(000).y', 'Employed_part-time_(000).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)_('000_Hours).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)_('000_Hours).y', 'Sex.x', 'State_and_territory_(STT):_ASGS_(2011).x',
## 'Employed_full-time_(000).x', 'Employed_part-time_(000).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)_('000_Hours).x',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)_('000_Hours).x', 'Sex.y', 'State_and_territory_(STT):_ASGS_(2011).x',
## 'Employed_full-time_(000).y', 'Employed_part-time_(000).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_full-
## time)_('000_Hours).y',
## 'Number_of_hours_actually_worked_in_all_jobs_(employed_part-
## time)_('000_Hours).y', 'Sex.x', 'State_and_territory_(STT):_ASGS_(2011).y' are duplicated
## in the result

```

```

# Remove columns where all values are NA
master_df <- Filter(function(x)!all(is.na(x)), master_df)

# Remove columns where NA values are more than 90% of records
master_df <- master_df[, which(colMeans(!is.na(master_df)) >= 0.1)] # filter the columns with more than

# Replace the remaining NAs with zero, if exist
nums <- unlist(lapply(master_df, is.numeric)) # select only numeric columns
for(i in 1:ncol(master_df[, nums])){
  return_numeric <- as.numeric(as.vector(as.numeric(unlist(master_df[, nums][, i])))) # converts import
  master_df[, nums][is.na(master_df[, nums][,i]), i] <- 0 # replace missing values with Zero
}

```

```

# Remove unnecessary characters from col names
names(master_df) <- sub("\\.x", "", names(master_df))

# Encode date column into Date format
master_df$Date <- as.Date(master_df$Date, format = "%Y-%m-%d")

# Change column names in a data frame
# format: "new column name to be assigned" = "previous column name"
master_df <- master_df %>%
  rename(
    Area = `Greater_capital_city_and_rest_of_state_(GCCSA):_ASGS_(2011)` ,
    Employed_full_time = `Employed_full-time_('000)` ,
    Employed_part_time = `Employed_part-time_('000)` ,
    Number_of_hours_worked_full_time = `Number_of_hours_actually_worked_in_all_jobs_(employed_full-time)` ,
    Number_of_hours_worked_part_time = `Number_of_hours_actually_worked_in_all_jobs_(employed_part-time)` )

# Derive a column from the master data frame to indicate relevance to tech sector
tech_value <- c("ICT", "Info", "Tech", "Engin", "IT", "Sci", "Math", "Mechanic", "Network") # Keywords
master_df$Tech <- grepl(paste(tech_value, collapse="|"), master_df$Occupation, ignore.case = TRUE)

# Print occupation types specified as being relevant to tech sector
occupation_tech <- unique(grep(paste(tech_value, collapse="|"),
                                 master_df$Occupation, value=TRUE, ignore.case = TRUE))

# Add column into master data frame indicating combined number of jobs and worked hour
master_df <- master_df %>%
  group_by(Date) %>%
  mutate(Employed_total = Employed_full_time + Employed_part_time,
         Number_of_hours_worked_avg = (Number_of_hours_worked_full_time + Number_of_hours_worked_part_time)/2)

# check for problems
assertthat::assert_that(nrow(problems(master_df)) == 0, # assert that there is NO problems
                        msg="There is still problem/s, which you need to fix first")

```

```
## [1] TRUE
```

```
# print data frame dimensions
cat("data dimensions are: ", dim(master_df))
```

```
## data dimensions are: 78637 11
```

```
# Export the completed records
write_csv(master_df, "data/master_df_all.csv")

# Derive a data frame only containing occupation types relevant to tech sector
tech_df <- master_df %>% filter(Tech == TRUE)

# Export the completed records
write_csv(tech_df, "data/master_df_tech.csv")

##### Part C #####

```

```

# Generate data frame arranged by area and occupation
mdata <- melt(master_df, id=c("Date", "Occupation", "Sex", "Area", "Tech"))
df_by_occupation_area <- dcast(mdata, Date + Occupation + Area + Tech ~ variable, mean)

# Generate data frame arranged by occupation and gender
df_by_occupation_gender <- dcast(mdata, Date + Occupation + Tech + Sex ~ variable, mean)

# Generate data frame arranged by occupation
df_by_occupation <- dcast(mdata, Date + Occupation + Tech ~ variable, mean)

# Generate newly arranged data frame only containing occupation types relevant to tech sector
mdata_tech <- master_df %>%
  filter(Tech == TRUE) %>%
  # mutate(Employed_total_tech = Employed_full_time + Employed_part_time,
  #        Number_of_hours_worked_avg_tech = (Number_of_hours_worked_full_time + Number_of_hours_worked_
  #        melt(id=c("Date", "Occupation", "Sex", "Area", "Tech")))
df_tech <- dcast(mdata_tech, Date + Sex + Area ~ variable, mean)

# Generate tech sector data frame arranged by gender
df_tech_by_area <- dcast(mdata_tech, Date + Sex ~ variable, mean)

# Generate tech sector data frame arranged by area
df_tech_by_gender <- dcast(mdata_tech, Date + Area ~ variable, mean)

# Generate tech sector data total
df_tech_total <- dcast(mdata_tech, Date ~ variable, mean)

# Generate newly rearranged data frame with occupation types other than tech sector
mdata_others <- master_df %>%
  filter(Tech == FALSE) %>%
  # mutate(Employed_total_others = Employed_full_time + Employed_part_time,
  #        Number_of_hours_worked_avg_others = (Number_of_hours_worked_full_time + Number_of_hours_worked_
  #        melt(id=c("Date", "Occupation", "Sex", "Area", "Tech")))
df_others <- dcast(mdata_others, Date + Sex + Area ~ variable, mean)

# Generate data frame arranged by gender
df_others_by_area <- dcast(mdata_others, Date + Sex ~ variable, mean)

# Generate data frame arranged by area
df_others_by_gender <- dcast(mdata_others, Date + Area ~ variable, mean)

# Generate data frame total
df_others_total <- dcast(mdata_others, Date ~ variable, mean)

# Print the date with the highest volume tech sector jobs
Highest_job_date_tech <- df_tech_total[which(df_tech_total$Employed_total == max(df_tech_total$Employed_), 
Highest_job_date_tech$Date

## [1] "2021-05-01"

# Print the date with the highest volume in jobs excluding tech sector
Highest_job_date_others <- df_others_total[which(df_others_total$Employed_total == max(df_others_total$Employed_), 
Highest_job_date_others$Date

```

```

## [1] "2021-11-01"

# Print the date with highest worked hours in tech sector
Highest_wokred_hour_tech <- df_tech_total[which(df_tech_total$Number_of_hours_worked_avg == max(df_tech_
Highest_wokred_hour_tech$Date

## [1] "2021-02-01"

# Print the date with highest worked hours in jobs excluding tech sector
Highest_wokred_hour_others <- df_others_total[which(df_others_total$Number_of_hours_worked_avg == max(df_othe
Highest_wokred_hour_others$Date

## [1] "2021-11-01"

# Generate data frame comparing job volume development between tech sector and others
df_volume_job_tech_others <- data.frame(Date = df_tech_total$Date,
                                         Employed_tech = df_tech_total$Employed_total,
                                         Employed_others = df_others_total$Employed_total
)

# Generate data frame comparing number of hour worked between tech sector and others
df_worked_hour_tech_others <- data.frame(Date = df_tech_total$Date,
                                             Number_of_hours_worked_tech = df_tech_total$Number_of_hours_work
                                             Number_of_hours_worked_others = df_others_total$Number_of_hours_
)

# Generate data frame comparing job volumes in each occupation
mdata_others_by_occupation <- df_by_occupation %>%
  filter(Tech == FALSE) %>%
  group_by(Date) %>%
  select(-Tech, -Employed_full_time, -Employed_part_time, -Number_of_hours_worked_full_time,
         -Number_of_hours_worked_part_time)
mdata_others_volume_by_occupation <- melt(mdata_others_by_occupation, id=c("Date", "Occupation"))
mdata_others_colume_by_occupation <- dcast(mdata_others_volume_by_occupation, Date ~ Occupation, mean)
names(mdata_others_colume_by_occupation) <- gsub(" ", "_", names(mdata_others_colume_by_occupation))
df_job_volume_development_tech_others <- data.frame(Date = df_volume_job_tech_others$Date,
                                                       Employed_tech = df_volume_job_tech_others$Employed_t
)
df_job_volume_development_tech_others <- merge(df_job_volume_development_tech_others,
                                                 mdata_others_colume_by_occupation)
names(df_job_volume_development_tech_others) <- gsubfn(".", list(" " = "_", "," = "_"),
                                                       names(df_job_volume_development_tech_others))

##### Part D #####
#print summary of data frame generated in the previous step
head(df_volume_job_tech_others)

```

	Date	Employed_tech	Employed_others
## 1	1984-11-01	6.766399	13.44001
## 2	1985-02-01	6.747070	13.48750
## 3	1985-05-01	6.906259	13.68875
## 4	1985-08-01	6.992410	13.65062
## 5	1985-11-01	7.203985	14.08972
## 6	1986-02-01	7.345076	14.00195

```
summary(df_volume_job_tech_others)

##          Date      Employed_tech      Employed_others
##  Min.    :1984-11-01   Min.    : 6.747   Min.    :13.44
##  1st Qu.:1994-02-01   1st Qu.: 8.174   1st Qu.:16.13
##  Median :2003-05-01   Median :11.069   Median :19.05
##  Mean   :2003-05-02   Mean    :11.893   Mean    :19.49
##  3rd Qu.:2012-08-01   3rd Qu.:14.952   3rd Qu.:22.60
##  Max.   :2021-11-01   Max.    :19.083   Max.    :26.31
```

```
df_volume_job_tech_others %>% tail
```

```
##          Date      Employed_tech      Employed_others
## 144 2020-08-01   17.83269      25.00101
## 145 2020-11-01   18.38723      25.46150
## 146 2021-02-01   19.05385      25.62662
## 147 2021-05-01   19.08332      26.09547
## 148 2021-08-01   18.78730      25.79299
## 149 2021-11-01   18.87154      26.31336
```

```
df_volume_job_tech_others %>% head
```

```
##          Date      Employed_tech      Employed_others
## 1 1984-11-01   6.766399      13.44001
## 2 1985-02-01   6.747070      13.48750
## 3 1985-05-01   6.906259      13.68875
## 4 1985-08-01   6.992410      13.65062
## 5 1985-11-01   7.203985      14.08972
## 6 1986-02-01   7.345076      14.00195
```

```
min(df_volume_job_tech_others[, 1], na.rm=T)
```

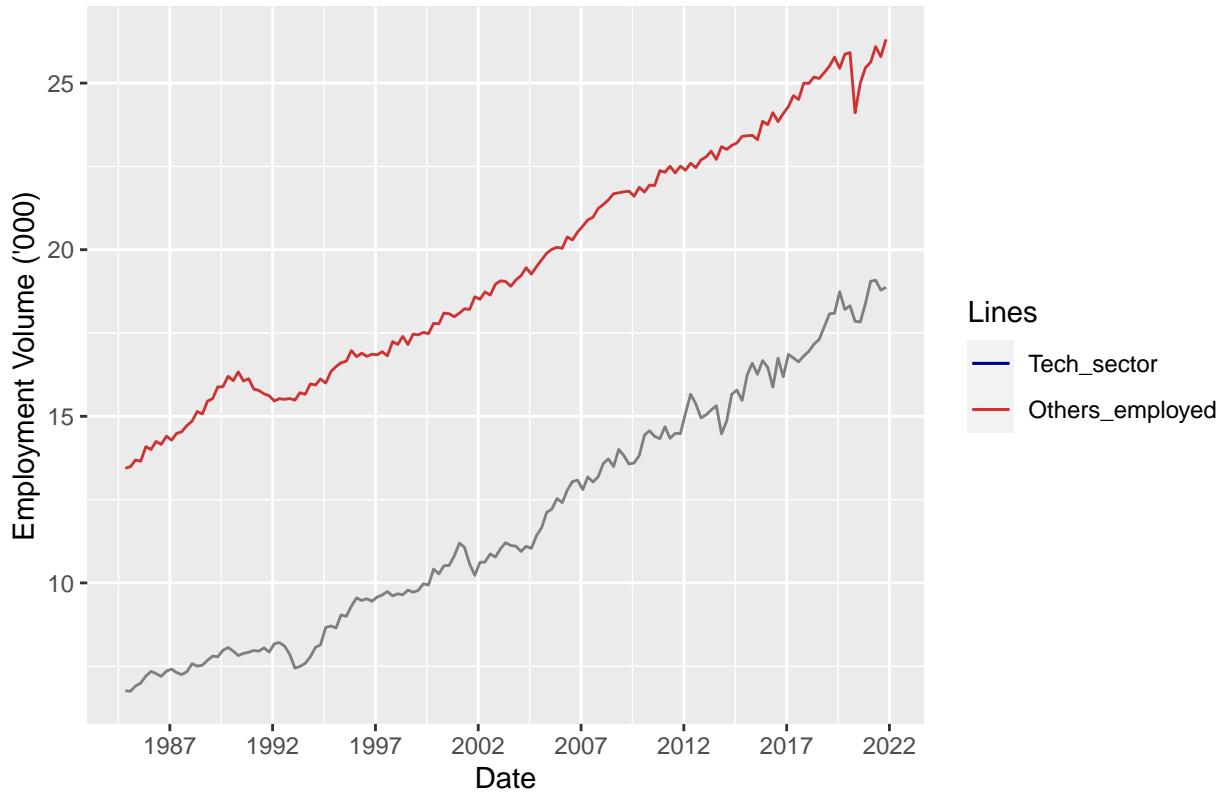
```
## [1] "1984-11-01"
```

```
max(df_volume_job_tech_others[, 1], na.rm=T)
```

```
## [1] "2021-11-01"
```

```
# Generate a graph showing development of volume of jobs in tech sector compared to others
cols <- c("Tech_sector"="darkblue","Others_employed"="brown3") # map the color in order to generate legend
graph_employment_tech_others <- df_volume_job_tech_others %>%
  ggplot(aes(x = Date)) +
  geom_line(aes(y = Employed_tech, color = "Tech_sector_employed")) +
  geom_line(aes(y = Employed_others, color = "Others_employed")) +
  scale_colour_manual(name="Lines",values=cols) +
  ggtitle("Employment Volume Tech Sector vs. Others") +
  scale_x_date(date_breaks = "5 year",
               labels=date_format("%Y")) +
  ggtitle("Employment Volume Tech Sector vs. Others") +
  ylab("Employment Volume ('000)")
print(graph_employment_tech_others)
```

Employment Volume Tech Sector vs. Others

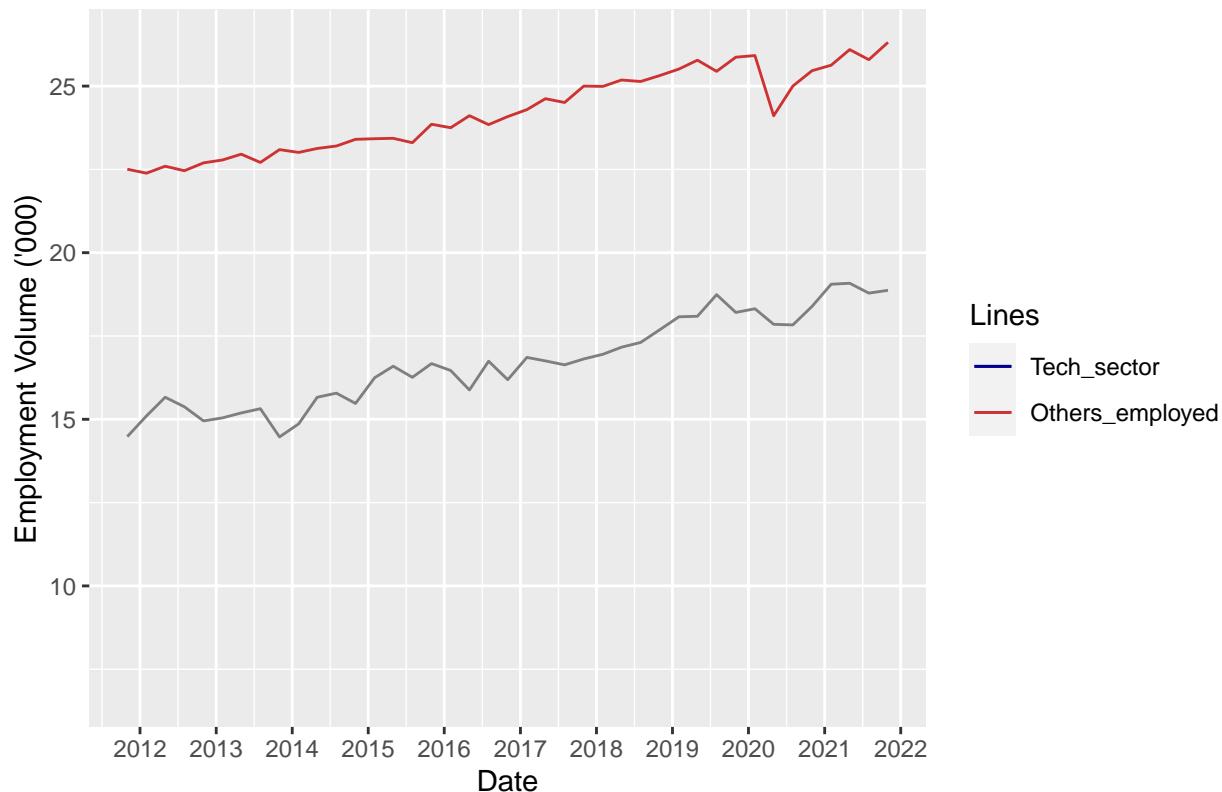


```
# limit x axis to recent 10 years
cols <- c("Tech_sector"="darkblue","Others_employed"="brown3") # map the color in order to generate legend
graph_employment_tech_others_xlimited <- df_volume_job_tech_others %>%
  ggplot(aes(x = Date)) +
  geom_line(aes(y = Employed_tech, color = "Tech_sector_employed")) +
  geom_line(aes(y = Employed_others, color = "Others_employed")) +
  scale_colour_manual(name="Lines",values=cols) +
  scale_x_date(date_breaks = "1 year",
               labels=date_format("%Y"),
               limits = as.Date(c(max(df_volume_job_tech_others>Date) - years(10),
                                 max(df_volume_job_tech_others>Date)))) +
  ggtitle("Employment Volume Tech Sector vs. Others (recent 10 years)") +
  ylab("Employment Volume ('000)")
print(graph_employment_tech_others_xlimited)
```

```
## Warning: Removed 108 row(s) containing missing values (geom_path).
```

```
## Warning: Removed 108 row(s) containing missing values (geom_path).
```

Employment Volume Tech Sector vs. Others (recent 10 years)



```
#####
# Part E #####
# For reproducibility
set.seed(123)

# randomly select 70% of the number of observations
index <- sample(1:nrow(df_job_volume_development_tech_others),
                 size = 0.7*nrow(df_job_volume_development_tech_others))

# subset other variables from the train data frame to include only the elements in the index
train <- df_job_volume_development_tech_others[index,]

# subset other variables from the test data frame to include only the elements in the index
test <- df_job_volume_development_tech_others[-index,]

nrow(train)

## [1] 104

nrow(test)

## [1] 45

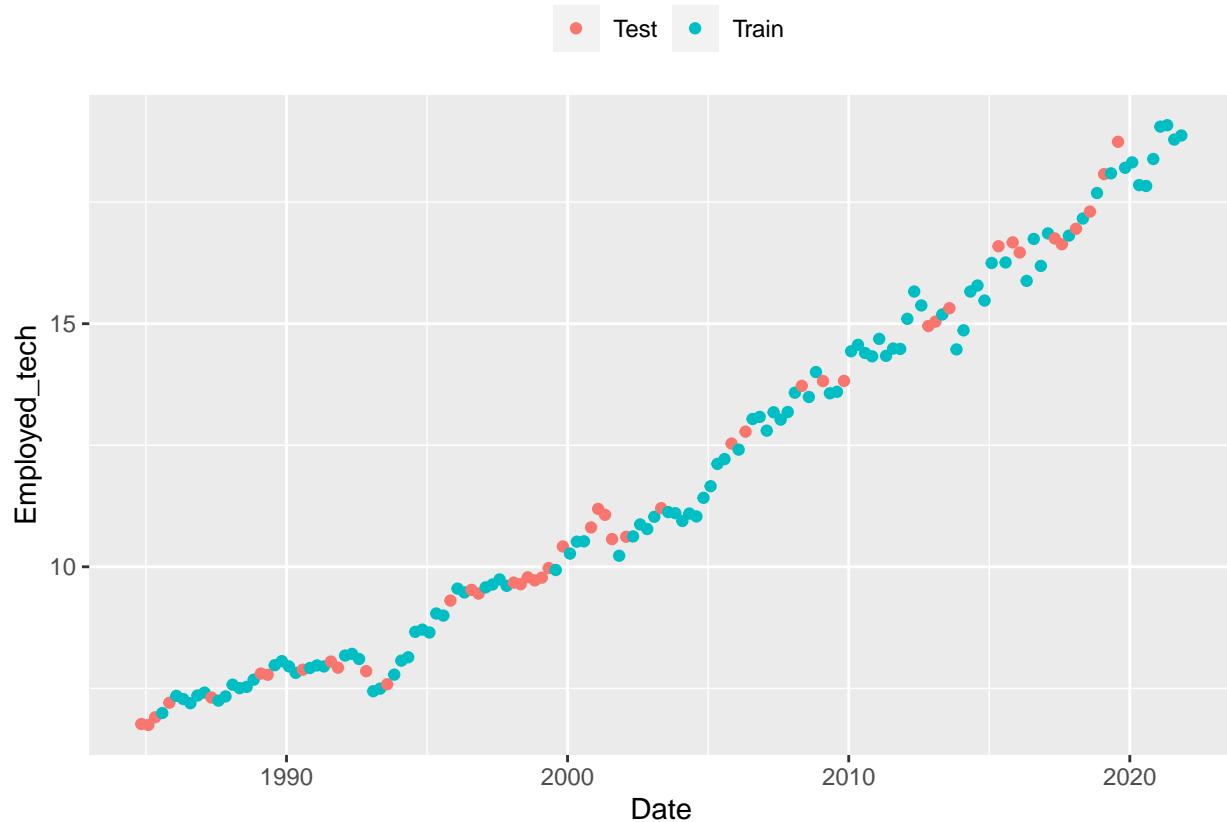
# Create a data frame with train and test indicator
group <- rep(NA,nrow(train)+nrow(test))
```

```

group <- ifelse(seq(1,nrow(train)+nrow(test)) %in% index, "Train", "Test")
df <- data.frame(Date=df_job_volume_development_tech_others$Date,
                  Employed_tech=df_job_volume_development_tech_others$Employed_tech,group)

# Plot test vs. train data frames
ggplot(df,aes(x = Date,y = Employed_tech, color = group)) + geom_point() +
  scale_color_discrete(name="") + theme(legend.position="top")

```



```

# Baseline model - predict the mean of the training data
best.guess <- mean(train$Employed_tech)

```

```

# Evaluate RMSE and MAE on the testing data
RMSE.baseline <- rmse(test$Employed_tech, best.guess)
RMSE.baseline

```

```
## [1] 3.623765
```

```

MAE.baseline <- mae(test$Employed_tech, best.guess)
MAE.baseline

```

```
## [1] 3.248703
```

```

#Multiple linear regression
# Create a multiple (log)linear regression model using the training data
lin.reg <- lm(log(Employed_tech+1) ~ ., data = train)

# Inspect the model
summary(lin.reg)

```

```

##
## Call:
## lm(formula = log(Employed_tech + 1) ~ ., data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -0.067223 -0.011122  0.000335  0.014383  0.046588 
## 
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|)    
## (Intercept)               1.452e+00  1.439e-01 10.087 3.03e-16  
## Date                     5.408e-05  8.240e-06  6.563 3.80e-09  
## Accommodation_and_Food_Services -9.539e-04  4.336e-04 -2.200  0.0305  
## Administrative_and_Support_Services 1.331e-03  5.915e-04  2.250  0.0270  
## Agriculture__Forestry_and_Fishing -9.523e-05  2.754e-04 -0.346  0.7303  
## Arts_and_Recreation_Services     1.523e-03  1.167e-03  1.305  0.1955  
## Construction                5.070e-04  2.603e-04  1.947  0.0548  
## Education_and_Training        -3.365e-04  3.740e-04 -0.900  0.3708  
## Financial_and_Insurance_Services 1.229e-03  5.110e-04  2.405  0.0183  
## Health_Care_and_Social_Assistance 1.631e-04  2.624e-04  0.622  0.5358  
## Manufacturing                3.343e-04  2.392e-04  1.398  0.1658  
## Mining                      2.827e-04  3.120e-04  0.906  0.3675  
## Other_Services                2.475e-04  4.784e-04  0.517  0.6062  
## Public_Administration_and_Safety -5.768e-04  5.074e-04 -1.137  0.2588  
## Rental__Hiring_and_Real_Estate_Services -9.633e-04  8.457e-04 -1.139  0.2578  
## Retail_Trade                 -1.997e-04  4.013e-04 -0.498  0.6200  
## Transport__Postal_and_Warehousing  1.048e-03  4.787e-04  2.190  0.0312  
## Wholesale_Trade              -9.004e-05  3.794e-04 -0.237  0.8130 
## 
## (Intercept)                   ***
## Date                         ***
## Accommodation_and_Food_Services *
## Administrative_and_Support_Services *
## Agriculture__Forestry_and_Fishing
## Arts_and_Recreation_Services
## Construction                  .
## Education_and_Training        *
## Financial_and_Insurance_Services *
## Health_Care_and_Social_Assistance
## Manufacturing
## Mining
## Other_Services
## Public_Administration_and_Safety
## Rental__Hiring_and_Real_Estate_Services
## Retail_Trade
## Transport__Postal_and_Warehousing    *

```

```

## Wholesale_Trade
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.02672 on 86 degrees of freedom
## Multiple R-squared:  0.9927, Adjusted R-squared:  0.9913
## F-statistic: 688.7 on 17 and 86 DF,  p-value: < 2.2e-16

# Multiplicative effect of "Mining" variable
exp(lin.reg$coefficients["Mining"])

## Mining
## 1.000283

# Apply the model to the testing data
test.pred.lin <- exp(predict(lin.reg,test))-1

# Evaluate the accuracy
RMSE.lin.reg <- rmse(test$Employed_tech, test.pred.lin)
RMSE.lin.reg

## [1] 0.3753509

MAE.lin.reg <- mae(test$Employed_tech, test.pred.lin)
MAE.lin.reg

## [1] 0.2830631

#Decision Tree
# rpart function applied to a numeric variable
rt <- rpart(Employed_tech ~ ., data=train)

fancyRpartPlot(rt)

# Predict and evaluate on the test set
test.pred.rtree <- predict(rt,test)

RMSE.rtree <- rmse(test$Employed_tech, test.pred.rtree)
RMSE.rtree

## [1] 0.7975804

MAE.rtree <- mae(test$Employed_tech, test.pred.rtree)
MAE.rtree

## [1] 0.6478226

# Cross-validation results (xerror)
printcp(rt)

```

```

## 
## Regression tree:
## rpart(formula = Employed_tech ~ ., data = train)
##
## Variables actually used in tree construction:
## [1] Date           Education_and_Training
##
## Root node error: 1399.5/104 = 13.457
##
## n= 104
##
##          CP nsplit rel error   xerror    xstd
## 1 0.781970      0 1.000000 1.031113 0.092787
## 2 0.093984      1  0.218030 0.238145 0.023953
## 3 0.075003      2  0.124046 0.161939 0.017625
## 4 0.014994      3  0.049043 0.091170 0.019745
## 5 0.010000      4  0.034049 0.083816 0.019797

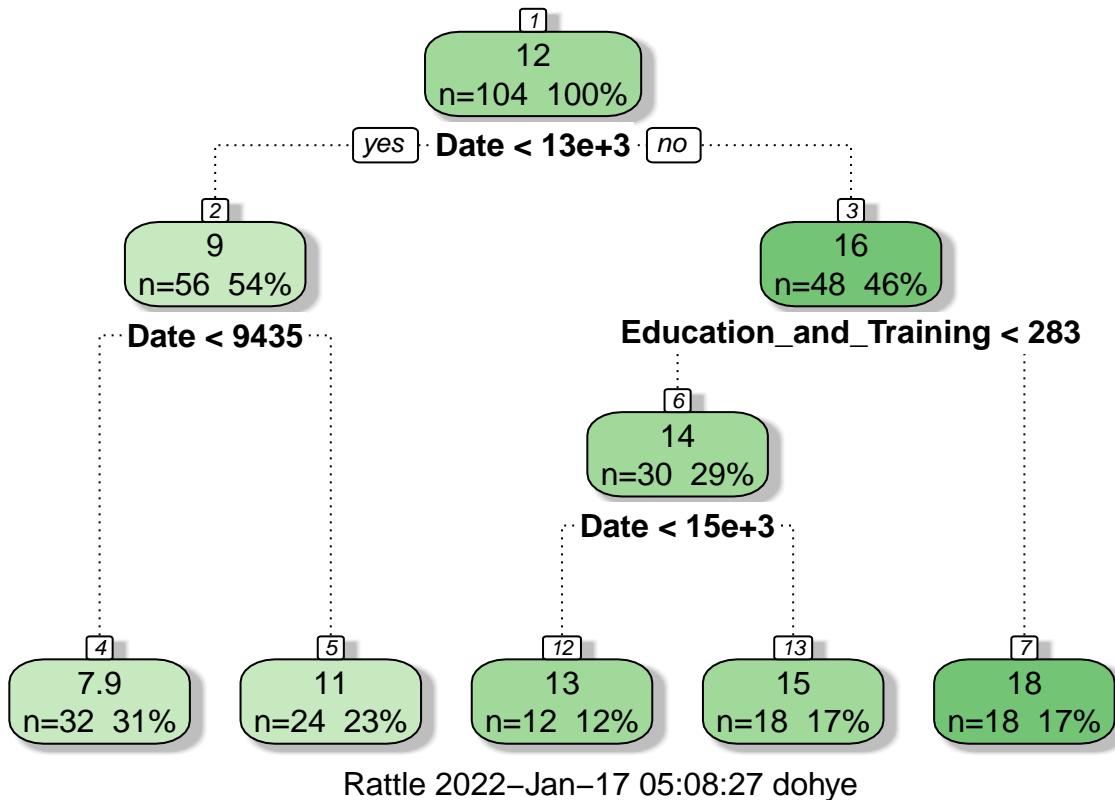
# Compute optimal CP
min.xerror <- rt$cptable[which.min(rt$cptable[, "xerror"]),"CP"]
min.xerror

## [1] 0.01

# Prune the tree
rt.pruned <- prune(rt,cp = min.xerror)

# Plot pruned tree
fancyRpartPlot(rt.pruned)

```



```

# Evaluate pruned tree on the test set
test.pred.rtree.p <- predict(rt.pruned,test)
RMSE.rtree.pruned <- rmse(test$Employed_tech, test.pred.rtree.p)
RMSE.rtree.pruned

## [1] 0.7975804

MAE.rtree.pruned <- mae(test$Employed_tech, test.pred.rtree.p)
MAE.rtree.pruned

## [1] 0.6478226

# Random forests
set.seed(123)

# Create a model with 1000 trees
rf <- randomForest(Employed_tech ~ ., data = train, importance = TRUE, ntree=1000)

importance(rf)

##                                     %IncMSE IncNodePurity
## Date                               19.269466   272.923392
## Accommodation_and_Food_Services    12.274719   11.650497
## Administrative_and_Support_Services 11.053875   16.062905
  
```

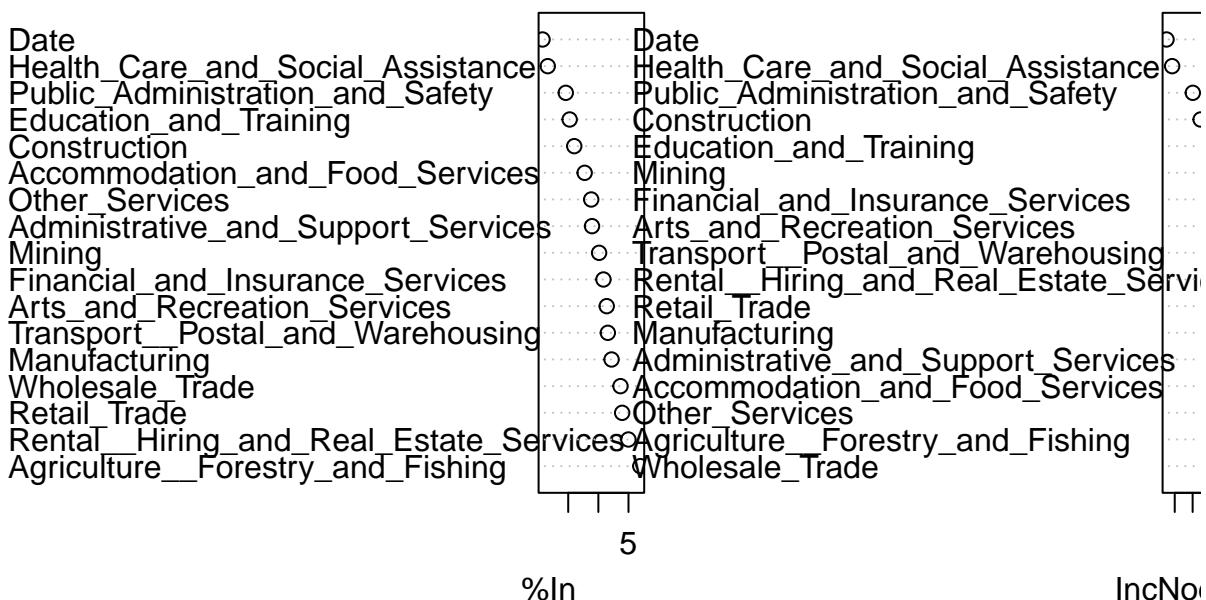
```

## Agriculture_Forestry_and_Fishing      3.040607    3.274280
## Arts_Recreation_Services             8.612227    56.883812
## Construction                        14.002774   177.765406
## Education_Training                  14.760791   110.250426
## Financial_Insurance_Services        9.156275    77.487982
## Health_Care_and_Social_Assistance   18.400583   257.211510
## Manufacturing                       7.808113    16.794758
## Mining                             9.875452    84.375893
## Other_Services                      11.187192   7.221025
## Public_Administration_and_Safety    15.407323   199.478738
## Rental_Hiring_and_Real_Estate_Servi 5.014103    20.224544
## Retail_Trade                        6.047478    19.547567
## Transport_Postal_and_Warehousing   8.427252    52.118162
## Wholesale_Trade                     6.327884    3.040388

```

```
varImpPlot(rf)
```

rf



```
# Compute optimal number of trees
which.min(rf$mse)
```

```
## [1] 44
```

```
# Calculate the importance of each variable
imp <- as.data.frame(sort(importance(rf)[,1],decreasing = TRUE),optional = T)
names(imp) <- "% Inc MSE"
imp
```

```

## % Inc MSE
## Date 19.269466
## Health_Care_and_Social_Assistance 18.400583
## Public_Administration_and_Safety 15.407323
## Education_and_Training 14.760791
## Construction 14.002774
## Accommodation_and_Food_Services 12.274719
## Other_Services 11.187192
## Administrative_and_Support_Services 11.053875
## Mining 9.875452
## Financial_and_Insurance_Services 9.156275
## Arts_and_Recreation_Services 8.612227
## Transport__Postal_and_Warehousing 8.427252
## Manufacturing 7.808113
## Wholesale_Trade 6.327884
## Retail_Trade 6.047478
## Rental__Hiring_and_Real_Estate_Services 5.014103
## Agriculture__Forestry_and_Fishing 3.040607

# Predict and evaluate on the test set
test.pred.forest <- predict(rf,test)
RMSE.forest <- rmse(test$Employed_tech, test.pred.forest)
RMSE.forest

## [1] 0.3410522

MAE.forest <- mae(test$Employed_tech, test.pred.forest)
MAE.forest

## [1] 0.2469949

# SVM
# trainx <- train[ , purrr::map_lgl(train, is.numeric)] # leave only numeric varialbes
# Train SVM model using train set
svm.model <- svm(Employed_tech ~ ., scale = T,
                  data=train,kernel="radial",cost=100,gamma=0.1);
# testx <- test[ , purrr::map_lgl(test, is.numeric)]

# Predict and evaluate on the test set
svm.pred <- predict(svm.model, test,decision.values =TRUE);
summary(svm.model)

## 
## Call:
## svm(formula = Employed_tech ~ ., data = train, kernel = "radial",
##      cost = 100, gamma = 0.1, scale = T)
## 
## Parameters:
##   SVM-Type:  eps-regression
##   SVM-Kernel: radial
##   cost: 100

```

```

##      gamma:  0.1
##      epsilon: 0.1
##
##
## Number of Support Vectors: 33

RMSE.svm <- rmse(test$Employed_tech, svm.pred)
MAE.svm <- mae(test$Employed_tech, svm.pred)

##### Part F #####
#Evaluate Results
# Create a data frame for error metrics of each method
accuracy <- data.frame(Method = c("Baseline", "Linear Regression", "Full tree", "Pruned tree",
                                    "Random forest", "Support Vector"),
                         RMSE = c(RMSE.baseline, RMSE.lin.reg,
                                   RMSE.rtree, RMSE.rtree.pruned, RMSE.forest, RMSE.svm),
                         MAE = c(MAE.baseline, MAE.lin.reg, MAE.rtree,
                                 MAE.rtree.pruned, MAE.forest, MAE.svm))

# Round the values and print the table
accuracy$RMSE <- round(accuracy$RMSE, 2)
accuracy$MAE <- round(accuracy$MAE, 2)

accuracy

##          Method RMSE   MAE
## 1        Baseline 3.62 3.25
## 2 Linear Regression 0.38 0.28
## 3       Full tree 0.80 0.65
## 4     Pruned tree 0.80 0.65
## 5    Random forest 0.34 0.25
## 6 Support Vector 0.48 0.35

#Print prediction
# Create a data frame with the predictions for each method
all.predictions <- data.frame(actual = test$Employed_tech,
                                baseline = best.guess,
                                linear.regression = test.pred.lin,
                                full.tree = test.pred.rtree,
                                pruned.tree = test.pred.rtree.p,
                                random.forest = test.pred.forest,
                                support.vector = svm.pred)

# First 10 observations of predictions
head(all.predictions, 10)

##      actual baseline linear.regression full.tree pruned.tree random.forest
## 1 6.766399 12.05043          6.490045  7.861442   7.861442    7.156777
## 2 6.747070 12.05043          6.708522  7.861442   7.861442    7.170190
## 3 6.906259 12.05043          6.807459  7.861442   7.861442    7.180887
## 5 7.203985 12.05043          6.870216  7.861442   7.861442    7.238689
## 11 7.310302 12.05043         7.280140  7.861442   7.861442    7.326189
## 18 7.803861 12.05043         7.765574  7.861442   7.861442    7.631973

```

```

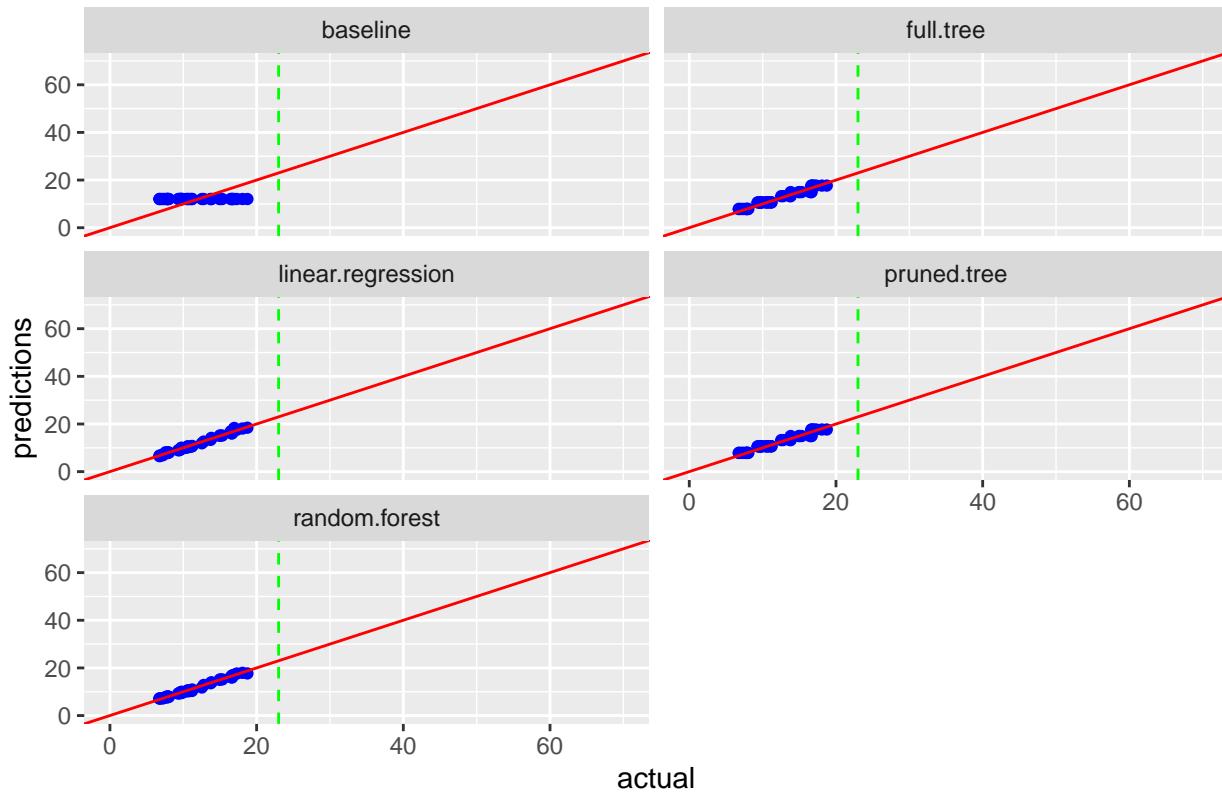
## 19 7.779840 12.05043      7.976532 7.861442 7.861442 7.898606
## 24 7.884218 12.05043      8.142751 7.861442 7.861442 8.055856
## 28 8.051184 12.05043      7.992188 7.861442 7.861442 7.870927
## 29 7.925592 12.05043      8.006559 7.861442 7.861442 7.945856
##   support.vector
## 1      8.197137
## 2      7.685155
## 3      7.545061
## 5      7.849466
## 11     7.285611
## 18     7.450947
## 19     8.036657
## 24     8.005247
## 28     7.949979
## 29     7.785498

# Convert the data frame in longer format
all.predictions <- gather(all.predictions, key = model, value = predictions, 2:6)

# Plot predicted vs. actual for each model
ggplot(data = all.predictions, aes(x = actual, y = predictions)) +
  geom_point(colour = "blue") +
  geom_abline(intercept = 0, slope = 1, colour = "red") +
  geom_vline(xintercept = 23, colour = "green", linetype = "dashed") +
  facet_wrap(~ model, ncol = 2) +
  coord_cartesian(xlim = c(0,70), ylim = c(0,70)) +
  ggtitle("Predicted vs. Actual, by model")

```

Predicted vs. Actual, by model



```
# Plot Predicted vs. actual by Random Forest model
random.forest.prediction <- data.frame(actual = test$Employed_tech,
                                         prediction = test$pred_forest)
random.forest.prediction
```

```
##           actual prediction
## 1     6.766399    7.156777
## 2     6.747070    7.170190
## 3     6.906259    7.180887
## 5     7.203985    7.238689
## 11    7.310302    7.326189
## 18    7.803861    7.631973
## 19    7.779840    7.898606
## 24    7.884218    8.055856
## 28    8.051184    7.870927
## 29    7.925592    7.945856
## 33    7.852994    7.698590
## 36    7.584705    7.828327
## 45    9.305477    9.236086
## 48    9.526157    9.401297
## 49    9.448940    9.467949
## 54    9.673935    9.508307
## 55    9.641958    9.634266
## 56    9.783893    9.429214
## 57    9.720148    9.940871
## 58    9.776988    9.931530
```

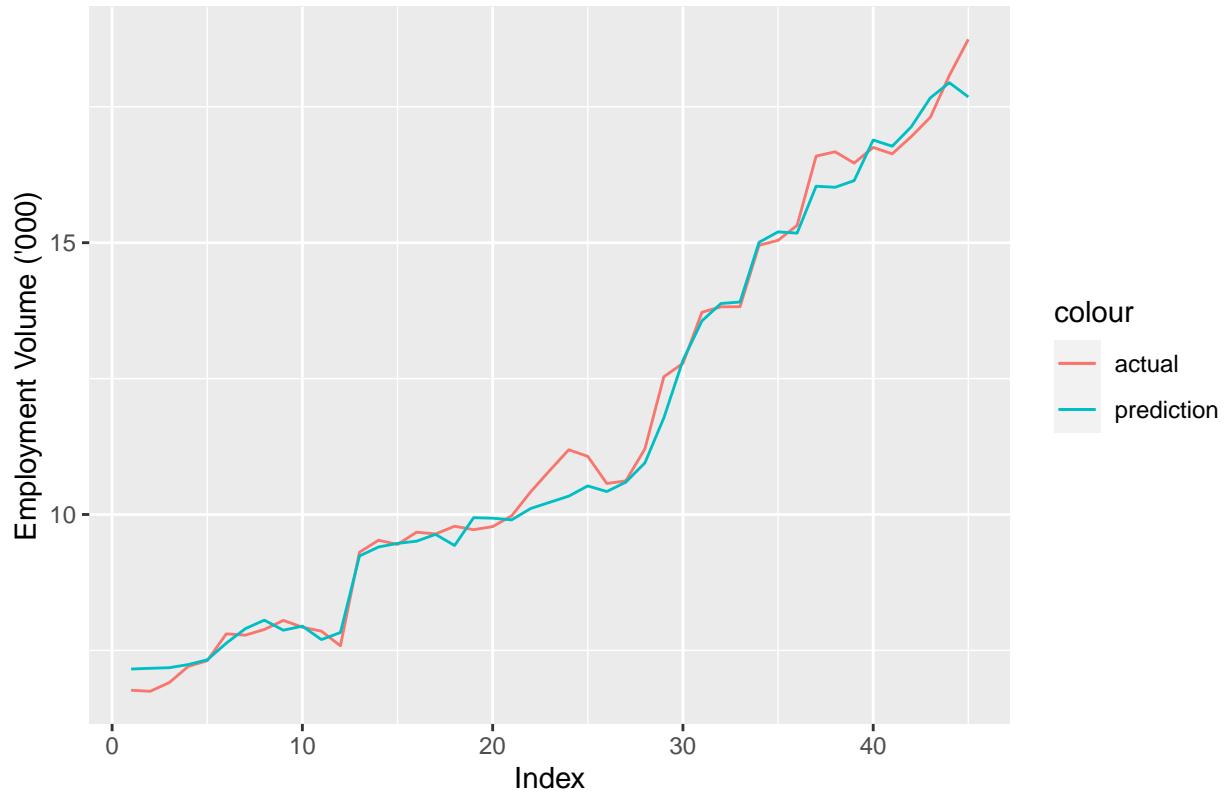
```

## 59   9.974754   9.901261
## 61  10.417863  10.110149
## 65  10.809358  10.224089
## 66  11.191156  10.337577
## 67  11.069500  10.525996
## 68  10.569790  10.421793
## 70  10.617940  10.595697
## 75  11.206245  10.948519
## 85  12.534807  11.779522
## 87  12.780181  12.836996
## 95  13.721785  13.561853
## 98  13.820091  13.881235
## 101 13.822742  13.908749
## 113 14.952276  15.007025
## 114 15.043909  15.199349
## 116 15.318442  15.176080
## 123 16.593138  16.038998
## 125 16.672569  16.019503
## 126 16.464713  16.142294
## 131 16.753101  16.888824
## 132 16.634246  16.774253
## 134 16.952214  17.132296
## 136 17.305651  17.663732
## 138 18.076146  17.944024
## 140 18.741404  17.681609

random.forest.prediction$index <- index(random.forest.prediction)
random.forest.prediction %>%
  ggplot(aes(x=index)) +
  geom_line(aes(y = actual, colour="actual")) +
  geom_line(aes(y = prediction, colour= "prediction")) +
  ylab("Employment Volume ('000)") +
  xlab("Index") +
  ggtitle(paste("Predicted vs. Actual by Random Forest"))

```

Predicted vs. Actual by Random Forest



```
#####
#set length of forecast (quarters)
duration <- 4

# Create a data frame used for forecasting
df <- df_job_volume_development_tech_others
df$date <- as.yearmon(df$date, "%Y-%m")

# Rearrange data frame depending on the prediction length
data_modifi <- df[,-which(names(df) == "Employed_tech")]
data_modifi$date <- as.yearmon(data_modifi$date) + (3/12)*duration
data_volume <- data.frame(date = df$date,
                           volume_employed = df[,which(names(df) == "Employed_tech")])
data <- merge(data_modifi, data_volume, all = TRUE)

# Drop dates out of range
data <- tail(data, -duration)

# leave only train data range
train <- data[as.yearmon(data$date)<=(as.yearmon(tail(data$date,1)) - (3/12)*duration),]

# subset elements in the index
test <- data[as.yearmon(data$date)>(as.yearmon(tail(data$date,1)) - (3/12)*duration),]

nrow(train)
```

```

## [1] 145

nrow(test)

## [1] 4

# Generate predictions using Random Forest
rf <- randomForest(volume_employed ~., data = train, importance = TRUE, ntree=1000)

# Compute optimal number of trees
ntree <- which.min(rf$mse)

# Regenerate predictions using optimal number of trees
rf <- randomForest(volume_employed ~., data = train, importance = TRUE, ntree=ntree)
test.pred <- predict(rf,test[,-which(names(test) == "volume_employed")])
test.pred

##      150      151      152      153
## 18.50815 18.63775 18.07068 18.63461

# Convert data frame into time series format
train.xts <- xts(train, order.by = as.Date(train$date))
test.xts <- xts(test, order.by = as.Date(test$date))

# Create a data frame indicating date range
pred2 <- rbind(train.xts[, "volume_employed"], test.xts[, "volume_employed"])

# Add corresponding date variables to predicted data
volume_employed_pred_xts <- xts(test.pred, order.by = as.Date(test$date))
names(volume_employed_pred_xts) <- "volume_employed_pred"
volume_employed_pred_xts

##           volume_employed_pred
## 2022-02-01          18.50815
## 2022-05-01          18.63775
## 2022-08-01          18.07068
## 2022-11-01          18.63461

# Create an data frame used for computing slope of change
df_slope <- data.frame(Date = index(volume_employed_pred_xts),
                        pred = volume_employed_pred_xts$volume_employed_pred)

# Compute slope of change
lm(df_slope$volume_employed_pred ~ index(df_slope))

## 
## Call:
## lm(formula = df_slope$volume_employed_pred ~ index(df_slope))
## 
## Coefficients:
## (Intercept) index(df_slope)
##           18.50971       -0.01877

```

```

# Compute average percentage change
APC <- ((df_slope[nrow(df_slope),2] - df_slope[1,2]) / df_slope[1,2]) * 100
APC

## [1] 0.6832955

# Comupte average percentage change per each quarter
APC_quarters <- APC/duration
APC_quarters

## [1] 0.1708239

# Merge predicted data with the original data
df_volume_pred_prev <- merge(pred2, volume_employed_pred_xts)
df_volume_pred_prev$volume_employed_pred[nrow(df_volume_pred_prev$volume_employed_pred)-duration] <-
  df_volume_pred_prev$volume_employed[nrow(df_volume_pred_prev$volume_employed_pred)-duration]

# Plot predicted values with the original data
df <- df_volume_pred_prev
graph_pred_prev <- plot(df[as.Date(as.yearmon(index(df)))>(as.Date(tail(as.yearmon(index(df)),1)-(3/12)),
                                         c("volume_employed", "volume_employed_pred")],
                           major.ticks = "months",
                           grid.ticks.on = "months",
                           grid.ticks.lty = 3,
                           main = paste(duration, "Quaters Forecast of Tech Sector Job Volume", sep=" "),
                           col = c("black", "blue"),
                           ylab = "Employment Volume in Tech Sector ('000)")

graph_pred_prev

```

4 Quarters Forecast of Tech Sector Job Volume

